

PEDRO NUNES PAGNUSSAT

+55 54 99960-0333 | np.pedro.np@gmail.com | linkedin.com/in/pedronp

SUMMARY

AI/ML Engineer with 5+ years building production systems in healthcare, telecom, and finance. Expertise in LLMs, RAG, and ML pipelines. Delivered **\$2.2M+ annual revenue** and **72x performance gains**. Proficient in Python, AWS/GCP, and end-to-end ML system design.

EXPERIENCE

Elsevier	Sep. 2025 – Present
Senior AI Engineer	Remote
<ul style="list-style-type: none">Architected multimodal RAG system for visual retrieval across healthcare resources, enabling medical professionals to search clinical images and tables alongside text.Designed OpenSearch indexing schema with custom chunking and metadata enrichment, improving search relevance for complex medical queries.Built high-performance image retrieval API with S3 and DynamoDB integration with async streaming, achieving sub-1-second response times at scale.Optimized table extraction pipeline, improving retrieval accuracy by 15% while reducing token consumption by 40%.	
Telus Digital	Nov. 2024 – Sep. 2025
Data Scientist	<i>Porto Alegre, Brazil (Hybrid)</i>
<ul style="list-style-type: none">Built ML-powered pricing optimization system analyzing product and market features to identify optimal price for Telecom Infrastructure, generating \$1.8M/year in incremental revenue.Developed automated international roaming analysis system covering 160+ countries (6x expansion from 25), replacing manual expert reviews and generating \$440K/year in revenue while freeing analysts for strategic work.	
Teia Labs (AI Startup)	Jan. 2024 – Nov. 2024
Junior Data Scientist	<i>Porto Alegre, Brazil (On-site)</i>
<ul style="list-style-type: none">Re-engineered embedding pipeline with parallelization and batching, achieving 72x speedup and enabling daily vector database refreshes.Built LLM-powered sentiment analysis pipeline processing 40,000+ companies for investment decision support.Deployed email summarization system using LLMs, reducing analyst review time by 4+ hours/week.Created AI agent research tool synthesizing web data in under 40 seconds.Integrated AI agents with RAG architecture across NoSQL and vector databases for unified cross-source retrieval.	
HP	Aug. 2021 – Jan. 2024
Research and Development Data Scientist	<i>Porto Alegre, Brazil (Hybrid)</i>
<ul style="list-style-type: none">Optimized SQL analytics queries reducing dashboard load times from minutes to under 5 seconds.Migrated 30+ OMEN gaming dashboards to Power BI, standardizing reporting infrastructure.Built a R&D for document intelligence: classification, NER-based auto-renaming, and semantic segmentation for document processing.	

EDUCATION

MBA in Project Management and Agile Methodologies	Dec. 2025 – Dec. 2026
<i>Pontifical Catholic University of Rio Grande do Sul (PUC-RS)</i>	<i>Porto Alegre, Brazil</i>
Bachelor of Science in Data Science and Artificial Intelligence	Mar. 2020 – Dec. 2024
<i>Pontifical Catholic University of Rio Grande do Sul (PUC-RS)</i>	<i>Porto Alegre, Brazil</i>

TECHNICAL SKILLS

ML/AI: Python, PyTorch, Scikit-learn, Hugging Face, LangChain, LangGraph OpenAI API, RAG, LLMs, NLP, Computer Vision
Data: SciPy, BigQuery, PostgreSQL, DynamoDB, MongoDB, OpenSearch, Vector Databases
Infrastructure: AWS (S3, Lambda, EC2), GCP, Docker, Airflow, FastAPI, Git, Linux
Spoken Languages: Portuguese (Native), English (C1), Spanish (Conversational), Chinese (HSK1)